



UNIMORE

UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA

UNIVERSITÀ DEGLI STUDI DI MODENA
E REGGIO EMILIA



An Overview of Generative Adversarial Networks (GAN)

Student

Cristian Cosci

Second Edition 2022

Contents

1	Introduction	4
1.1	GAN: Generator and Discriminator	6
1.1.1	Generator	6
1.1.2	Discriminator	7
1.2	Training Phase	8
2	Some GAN's variants	13
2.1	Deep Convolutional GAN (DCGAN)	13
2.2	Conditional GAN (cGAN)	15
2.2.1	cGAN's Generator	16
2.2.2	cGAN's Discriminator	17
2.3	Conditional Deep Convolutional GAN (cDCGAN)	17
3	Implementation and results	18
3.1	GAN	18
3.1.1	Generator	18
3.1.2	Discriminator	19
3.1.3	Training and results	19
3.2	cGAN	21
3.2.1	Generator	21
3.2.2	Discriminator	21
3.2.3	Training and Results	21
3.3	DCGAN	23
3.3.1	Generator	23
3.3.2	Discriminator	24

3.3.3	Training and results	24
3.4	cDCGAN	26
4	Conclusion	27

Chapter 1

Introduction

Generative Adversarial Networks were originally introduced by Goodfellow et al. [3] in 2014. He introduced GAN as a new framework for estimating generative models via an adversarial process, in which a generative model G captures the data distribution, while a discriminative model D estimates if the sample came from the training data rather than G .

In other word, GAN are a particular neural network that take random noise as input and generate picture as outputs. This picture appear to be a sample from the distribution of the training set (e.g. MNIST or other image based dataset).

A GAN structure consist in two different (and with different scope) models which are trained simultaneously and in an adversarial way:

- A **Generative** model (also called **Generator**) that captures the distribution of training set.
- A **Discriminative** model (also called **Discriminator**) that estimates the probability that a sample is a real or a fake picture (e.g. is from training data or is generated by the other model).

The simplest explanation about GANs training is to think at Generator as a banknotes' counterfeiter and at Discriminator as a policeman specialized in banknote recognition (see Figure 1.1).

The policeman are taught how to determine if a dollar bill is either genuine or fake. Samples of real dollar bills from the bank and fake money from the counterfeiter are used to train the policeman. However, from time to time, the counterfeiter will

attempt to pretend that he printed real dollar bills. Initially, the police will not be fooled and will tell the counterfeiter why the money is fake. Taking into consideration this feedback, the counterfeiter hones his skills again and attempts to produce new fake dollar bills. As expected the policeman will be able to both spot the money as fake and justify why the dollar bills are fake [5].

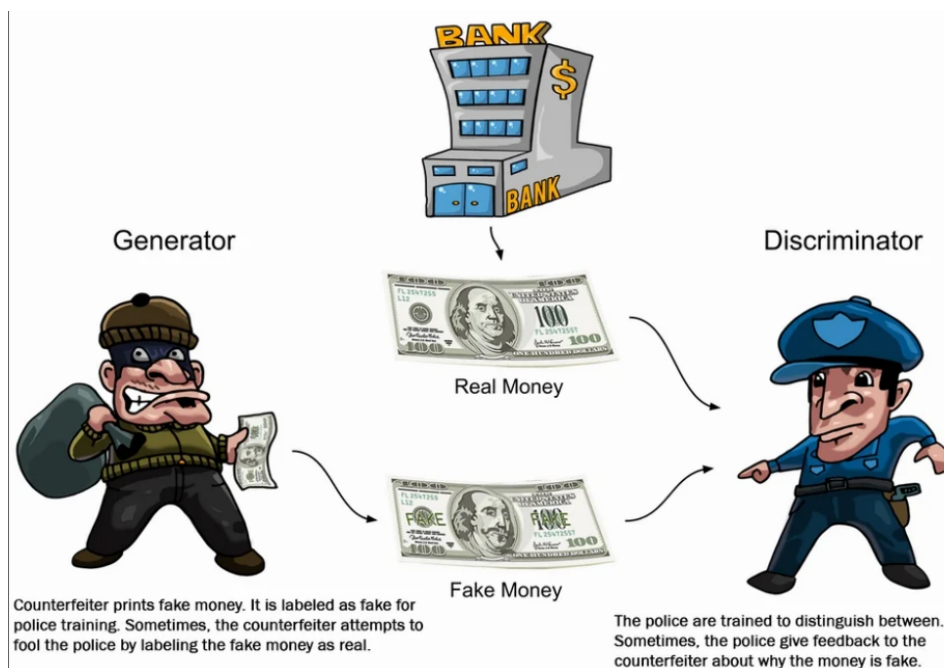


Figure 1.1: Generator vs Discriminator examples [5].

In other words the Generator is trying to learn the distribution of real data (by taking random noise as input, and producing realistic-looking images) and is the network which we're usually interested in. On the other hand, the Discriminator tries to classify whether the sample has come from the real dataset, or is fake (generated by the generator).

During the game the goal of the generator is to trick the discriminator into "thinking" that the data it generates is real. The goal of the discriminator, on the other hand, is to correctly discriminate between the generated (fake) images and real images coming from some dataset (e.g. MNIST).

Therefore, in the training phase the Discriminator model give a feedback to the Generator model in order to tell him if his work is good or bad. In this manner the Generator correct and improve his generation phase. However, at the same time, also

the Discriminator improve his classification while taking as input also real sample from the dataset and receiving a supervised feedback on his work.

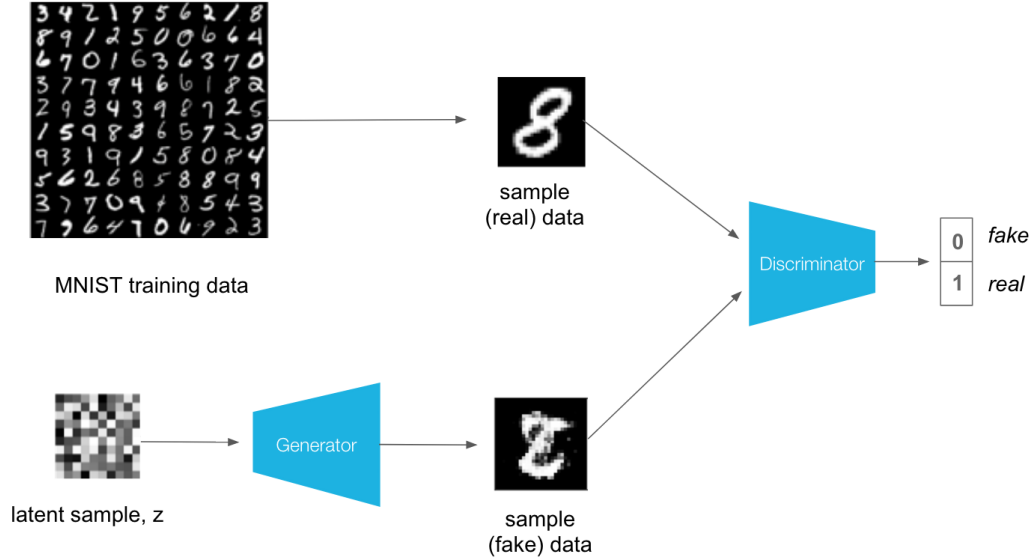


Figure 1.2: Simple GAN structure [1].

1.1 GAN: Generator and Discriminator

In this section i provide a little bit more deep explanation about the principal components of GAN. I will describe the structure of Generator and Discriminator, the training phase and the objective function that guides it.

1.1.1 Generator

Generator is a Neural Network, which given a dataset X_{real} tries to capture his distribution, by producing images X_{fake} from noise Z as input. Noise input is a random set of values, usually sampled from a multivariate-gaussian distribution. This is often called **latent vector** and that vector space is called **latent space**. The GAN's Generator acts quite like the decoder component of VAE: it project latent space to an image. The difference is that the Generator's latent space is not forced to learn exactly a Gaussian distribution but it can learn and model more complex distribution.

A problem that can bring GAN to poor performance is **Mode collapse**, which happens when the Generator can only produce a single type of output or a small set of outputs. This may happen due to problems in training, such as the generator finds a type of data that is easily able to fool the Discriminator and thus keeps generating that one type.




Figure 1.3: GAN's Generator example [7].

1.1.2 Discriminator

The Discriminator is more simple to understand than the Generator. The Discriminator is like a model trained to classify if an input image is real or is generated by another model. In other words the Discriminator is a binary classifier trained with a supervised classification mode.

The Discriminator task is to predict a label (e.g. *True* or *Fake* or 0-1 label) from an input. While doing so, the supervised train gives to it a feedback and this permit the model to learn a set of parameters (weights and bias) in order to map the correct label.

The principal scope of the Discriminator is to give a feedback to the Generator model, and to improve its generation phase. Therefore the Discriminator is used to train the Generator and after the training process, the Discriminator model is discarded as we

are interested in the generator.

Sometimes, the generator can be repurposed as it has learned to effectively extract features from examples in the problem domain. Some or all of the feature extraction layers can be used in transfer learning applications using the same or similar input data.




Figure 1.4: GAN's Discriminator example [7].

1.2 Training Phase

In this section i will describe the principal aspects of training phase on Generative Adversarial Networks.

But first a brief recap: the two models, the Generator and Discriminator, are trained together. The generator generates some samples, and these, along with real examples from the training dataset, are provided to the discriminator which classify these as real or fake. The Discriminator receive a feedback (is a supervised task) and is then updated to get better performance at discriminating real and fake samples in the next round. Also the Generator is updated based on how well, or not, the generated samples fooled the Discriminator.

Before start the explanation, let's denote some of the key elements:

- $X \rightarrow$ Training samples
- $Z \rightarrow$ Noise or Latent Vector
- $D \rightarrow$ Discriminator
- $G \rightarrow$ Generator
- $X_{real} \rightarrow$ a sample from training dataset
- X_{fake} or $G(z) \rightarrow$ Generator's output
- $D(x) \in (0, 1) \rightarrow$ Discriminator's output

The training of the two models (D and G) is done in an alternating way:

1. **Step 1:**

- The Generator take Z as input and produce $G(x)$ (X_{fake}).
- X_{fake} and a sample X_{real} are passed to the Discriminator.
- The Discriminator gives $D(x)$ as output. This tells the probability score of each input passed to D to be real or fake.
- As a normal supervised training, the predictions are compared with the ground truth, and a **Loss** is calculate (e.g. Binary Cross-Entropy).
- The Gradient is the backpropagated only through the Discriminator and is parameters are updated.

2. **Step 2:**

- The Generator produce an image $G(x)$, that is again passed through the Discriminator.
- The Discriminator gives a prediction $D(x)$ as before, but only for image passed by the Generator.
- The Loss is computed as before.
- In this step, because we want to enforce the Generator to produce images, as similar to the real images as possible (i.e., close to the true distribution), the Loss is backpropagated only through the Generator network, and its parameters are optimized suitably.

It is clear that the Discriminator is very important for the Generator's training. The Discriminator permits to compute a Loss and that is backpropagate through the generator to improve and update its weights and bias.

However there is a need for both networks to be strong enough because if the Discriminator is a weak classifier, also low quality images produced by Generator will be accepted as real images and then there will be no improvement for the Generator's images production. In other hand, if the Generator is weak, it will not be able to fool the Discriminator.

The best case scenario is when the Discriminator is no more able to distinguish if an image produced by the Generator is real or fake, so its predictions are like a toss of a coin i.e. it choose randomly (50% is fake and 50% is real) (it is also called Nash equilibrium). At the same time, obviously, the Generator has to be strong enough to produce very high quality images that his consistent with data distribution.

As described, both the Generator and the Discriminator's training is based on the classification score given as output by the last layer of D model. The output score is based on a binary classification problem, so the Binary Cross-Entropy function is used to compute the Loss:

$$L(\hat{y}, y) = -\frac{1}{N} \sum_{i=1}^N [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

We can go a little bit more deep on that Loss function:

- The negative sign at the beginning of the equation is there because the output of a neural network is normalized between 0 and 1, compute the \log of a number in that range result in a value less than zero.
- The $\frac{1}{N}$ term identify the Loss' means computed across the batch (N refer to batch size).
- \hat{y}_i is the Discriminator's prediction.
- y_i is the true label.
- The equation's first term is valid when the true label is 1 (real), and the second is valid when the true label is 0 (fake).

As said before, we can describe GAN’s training phase as a **minmax** game played by Generator and Discriminator. Is possible to explain it with the following value function $V(G, D)$ [3]:

$$\min_G \max_D V(D, G) = E_{x \sim p_{data}(x)}[\log D(x)] + E_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

However, the previously value function may not provide sufficient gradient for the generator to learn the data distribution so well [3]. Training this way will achieve only half the objective. Though the discriminator definitely becomes more powerful for it can now easily discriminate the real from the fake, the generator lags behind. It has still not learned to produce realistic-looking images. The consequence is that quickly in the learning, when the Generator’s performance is poor, the Discriminator can distinguish real and fake samples with high confidence. In this case, $-\log(1 - D(G(z)))$ saturates. Hence, rather than training G to minimize $-\log(1 - D(G(z)))$, better performances occurs training G to maximize $-\log D(G(z))$ [3].

To get more detail, we can start from the Discriminator’s **objective function**. The Discriminator is a binary classifier, that given an input x , returns a probability between 0 and 1 called $D(x)$. For instance the label for X_{real} is 1 and the label for X_{fake} is 0. A value for $D(x)$ closer to 1 means that the discriminator predicts that the input is a real image. On the other hand, a value for $D(x)$ closer to 0 means that the input is fake.

Thus, the objective function of the Discriminator becomes:

- **Maximizing** $D(X_{real})$ and **Minimizing** $D(X_{fake})$, where $X_{fake} = G(Z)$.

The objective function is modelled by Binary Cross-Entropy loss function, and we get the final formula substituting terms in the loss equation [7]:

- $y = 1$ and $\hat{y} = D(X_{real})$

$$D_{loss_{real}} = -\log D(X_{real}) \tag{1.1}$$

- true label y is 0, and the predicted output \hat{y} is $D(X_{fake})$ where X_{fake} is equal

to $G(z)$. Putting these values in the BCE loss function, we get:

$$D_{loss_{fake}} = -\log(1 - D(G(z))) \quad (1.2)$$

- Therefore, the cumulative Discriminator loss becomes:

$$D_{loss} = D_{loss_{real}} + D_{loss_{fake}} \quad (1.3)$$

$$D_{loss} = -\log D(X_{real}) - \log(1 - D(G(z))) \quad (1.4)$$

The Generator's **objective function** has the scope to Maximizing $D(G(Z))$ i.e. bringing it closer to 1 and fool the Discriminator. For this objective, i.e., to maximize the probability $D(G(z))$ by the Discriminator, the true label y is 1, and the predicted output \hat{y} is $D(G(z))$. Putting these values in the BCE loss function, we get [7]:

$$G_{loss} = -\log D(G(z)) \quad (1.5)$$

These are the most important thing to know about GANs, their components and their training phase. In the next Chapters i will describe the various type of GANs which were implemented after the Goodfellow et al.'s paper and their implementations.

Chapter 2

Some GAN's variants

After Goodfellow's article on GAN [3], many researchers have become interested in the Generative Adversarial Networks. Many variants of the original GAN architecture have been proposed and studied.

In the next section i will propose some of the most important variations like DCGAN, cGAN and cDCGAN.

2.1 Deep Convolutional GAN (DCGAN)

Deep Convolutional Generative Adversarial Network, also know as DCGAN is a particular and new GAN's architecture that improve his quality using **convolutional layers**. It was introduced around 2016 by Alec Radford and other researcher in a paper published at ICLR conference [6].

The main components introduced in DCGAN is:

- **fractionally-strided convolutional layers** in the Generator, with the scope of upsample the images (see Figure 2.1).
- **strided convolutional layers** in the Discriminator, with the scope of down-sample the images.




Figure 2.1: DCGAN generator used in the Radford’s introduction paper for scene modeling. A 100 dimensional uniform distribution Z is projected to a small spatial extent convolutional representation with many feature maps. A series of four fractionally-strided convolutions (in some recent papers, these are wrongly called de-convolutions) then convert this high level representation into a 64×64 pixel image. Notably, no fully connected or pooling layers are used [6].

In the Figure 2.2 we can see an example of a strided convolutional layer. In DCGAN the authors used a Stride of 2, meaning the filter slides thorough the image moving 2 pixels per step. The use of strided convolutional layer replaced the use of pooling operation like **maxPooling** or **avgPooling** because these operation have not learn-able parameter. The reason why they used a strided convolutional layer is because in this manner the network (the Discriminator in that case) is allowed to learn its own spatial downsampling [6].

The **Fractionally-strided convolution** operation is in essence the opposite of a convolution operations. In a traditional convolution operation there is a downsampling of the input dimension in the output (see Figure 2.2). Instead, in a fractionally-strided convolution operation, is produced a larger output from the input.

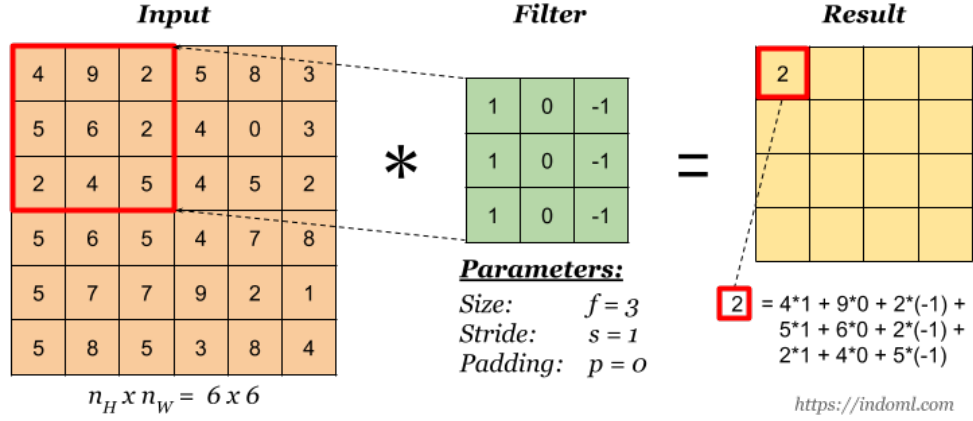


Figure 2.2: Example of a traditional convolution operation [2].

The power of this upsampling operation is that fractionally-strided convolution has **learnable parameters**. This parameters can be trained in order to follow the objective function and maximize the performance.

2.2 Conditional GAN (cGAN)

In the previous section i explained DCGAN, that is a particular GAN which differs from the original version by using particular convolutional layer (both in the Generator and Discriminator). The main functionality once trained is exactly the same as GAN: take as input an noise vector from latent space and give as output an image that match the training dataset distribution.

Traditional GAN is trained in a completely unsupervised and unconditional fashion, meaning no labels involved in the training process. Though the GAN model is capable of generating new realistic samples for a particular dataset, we have zero control over the type of images that are generated [8].

This particular variant, **Conditional GAN**, were introduced by Mehdi Mirza and Simon Osindero [4]. To understand the cGAN mechanism, think about a traditional GAN trained to generate images that match a dataset with handwritten number (e.g. MNIST dataset). With a GAN you pass a noise vector as input and you get an images as output. You don't have power on the number that GAN will output to you (e.g. it may be a 9 or it may be a 1). Of course you can have explore the latent space and find the area of the distribution which you can sample a latent vector that

will give you a specific number as output, but to obtain this you have to do a type of brute-force until you find the area which corresponds to a certain number. This problem occurs with any dataset you used to train a GAN. The cGAN is a solution to this by giving some type of control on generation phase. In the Figure 2.3 we can see the differences on the structure of GAN and cGAN.

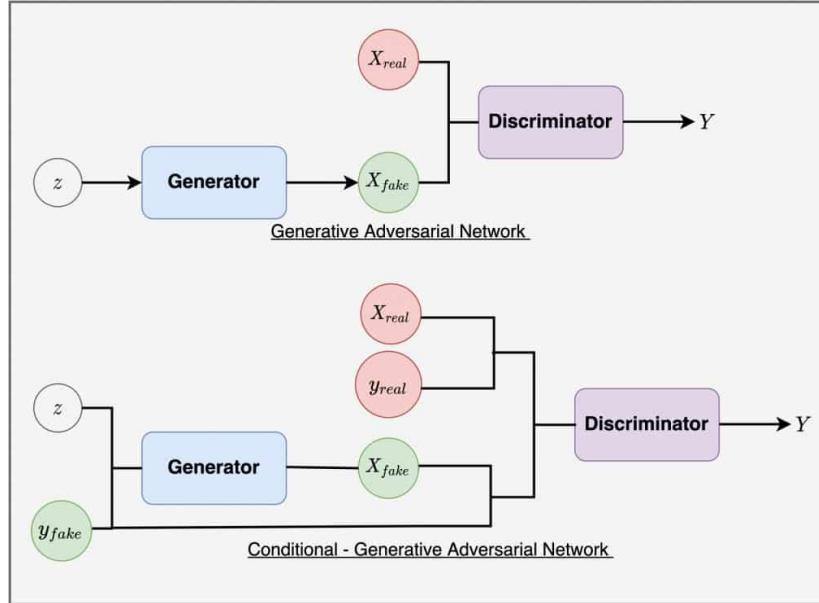


Figure 2.3: Flow Diagram representing GAN and Conditional GAN [8].

As you can see in the Figure 2.3, the Conditional GAN structure introduce an external information (all the other components is exactly the same as original GAN). This information could be a class label and is used to condition the model during the training so that is possible to condition both the generation phase.

2.2.1 cGAN's Generator

In GAN, the Generator use only a latent vector of noise to generate a sample image. In a conditional GAN the Generator also has an auxiliary information that tells which particular class sample to produce. So, there is a **conditioning label y** that is passed to G with noise vector Z :

$$GAN \rightarrow G(Z) \quad (2.1)$$

$$cGAN \rightarrow G(Z, y) = Z|y \quad (2.2)$$

Therefore, is not enough for the Generator to produce a realistic-looking images, but it is equally important that the generated sample match the label y . If the Generator has been properly trained, is possible to condition the image generation using a label.

2.2.2 cGAN's Discriminator

At the same time, the Discriminator receives both true and false examples with labels. In this way, the discriminator learns both to recognize real data from false ones and whether these samples match the labels. Then, the discriminator will reject realistic looking samples (generated by Generator) but not matching the label in the same way as fake samples.

2.3 Conditional Deep Convolutional GAN (cDCGAN)

Conditional Deep Convolutional GAN is a conditional GAN that use the same convolution layers as DCGAN that is described previously. cDCGAN generate more realistic images than cGAN thanks to convolutional layers and in the next Chapter some examples will be presented.

Chapter 3

Implementation and results

In order to see the differences between the models, in terms of generation performance, i have developed all the models previously described.

It is all available at this **GitHub repo** [Generative_Adversarial_Networks__Overview](#). For the execution's information about the project and different models you can see the *README* file in the GitHub repo.

In the next section i will describe the models structure and i will show some generation results. It starts with vanilla GAN up to Conditional Deep Convolutional GAN.

3.1 GAN

As described in the previous sections, the GAN's Discriminator and Generator has the simplest structures with respect to models like DCGAN. Both Discriminat and Generator are definided as an MLP.

3.1.1 Generator

The Generator has 5 layers with input and output layers included.

The first layer input size correspond to the latent vector size (100) and then i proceded doubling the dimension start from 128 until 1024 in the second-last layers. In the end the output layer size is the same as MNIST dataset's image size (i.e. 784 given that the image is 28 x 28). I used LeakyReLU as activations function in the intermediate

levels and a Batch Normalization. Leaky ReLU is a type of activation function that helps to prevent the function from becoming saturated at 0. It has a small slope instead of the standard ReLU which has an infinite slope.

Leaky ReLU is a modification of the ReLU activation function. It has the same form as the ReLU, but it will leak some positive values to 0 if they are close enough to zero. It uses leaky values to avoid dividing by zero when the input value is negative, which can happen with standard ReLU when training neural networks with gradient descent.

Batch normalization (also known as batch norm) is a method used to make training of artificial neural networks faster and more stable through normalization of the layers' inputs by re-centering and re-scaling.

The activation function used in the output layer is **tanh** to ensure that the pixel values are mapped in line with its own output, i.e., between (-1, 1).

3.1.2 Discriminator

The Discriminator model is defined as a binary classifier. The input size matches the flattened image size from dataset or Generator. The i reduced the size over the 3 layers until the output layer which has a size of 1 (i.e. the probability of a binary problem -> real or fake). The final layer has the sigmoid activation function, which squashes the output value between 0 (fake) and 1 (real).

3.1.3 Training and results

For the training process i decided to do a standard and simple process (considering the models structure). I trained the GAN using 200 epochs and then see if there were some improvements on generation process. As i described before the GAN training phase i too complex and is not sufficient to whatch te Loss to have an explaining about model performances. Of course, the best case scenario is when the Discriminator do not recognize anymore if the input is real or fake and the Generator give as output realistic-looking image (Nash equilibrium), but it is very difficult for that to happen. So, in the Table 3.1 is possible to see the generation improvement during the GAN's training phase: the images are initially very blurry and some figures does not make sense, instead in the last epoch, they are very sharp and realistic.



(a) Epoch 1



(b) Epoch 70



(c) Epoch 130



(d) Epoch 200

Table 3.1: Performance improvement during training in terms of generations.

3.2 cGAN

The Conditional Gan model is essentially the same as Vannilla Gan implementation, the only differences is that we need to give as input to the model also a vector which has the task to condition both the Discriminator and Generator. This vector is needed to guide the Generator to produce samples that match the class we use to condition it. In the case of Discriminator, the conditional vector is used to make the model capable to recognize if the sample given as input match the class.

3.2.1 Generator


The cGAN's Generator model structure has the same number of layers with the same activation functions and number of neurons for each level as Vanilla GAN. The difference is in the input layer, where the input size is not equal to a noise vector from latent space (e.g. size of 100) but is equal to this vector plus another vector called **conditional vector**. The size of this vector depends from the encoding we want to use and on the number of classes. For example, since i want to use cGAN to generate MNIST handwritten digits, i used a one-hot-econding for the 10 classes. So, in that case, the input size of the input layer is $100 + 10$ (noise vector dim + conditional vector dim).

3.2.2 Discriminator


In the Discriminator we have the same situation, the only difference is in the input layer's input size, where, instead of take only the flattened image as input, the model take also the conditional vector (e.g. in that case $784+10$).

3.2.3 Training and Results

For the training process i decided to do a standard and simple process (considering the models structure). I trained the cGAN using 150 epochs and then see if there were some improvements on generation process. In the Table 3.2 is possible to see the generation improvement during the cGAN's training phase: the images are initially very blurry and some figures does not make sense, instead in the last epoch, they are very sharp and realistic.



(a) Epoch 1



(b) Epoch 50



(c) Epoch 100



(d) Epoch 150

Table 3.2: Performance improvement during training in terms of generations.

3.3 DCGAN

Deep Convolutional Generative Adversarial Network also known as DCGAN, improves the quality of GANs using convolutional layers. In this implementation i preferred using a different dataset, in order to see the better generation power of DGGAN with respect to GAN. I used the Anime Face Dataset that consists of 63,632 high-quality anime faces, which were scraped from getchu, then cropped using the anime face-detection algorithm. Once DCGAN is trained, the Generator will be able to produce realistic-looking anime faces, like the ones shown in Figure 3.1.



Figure 3.1: Example images from the Anime Face Dataset [9].

3.3.1 Generator

As described before, the Generator model in DCGAN uses a particular type of Convolutional Layer called **Convolution 2D Transpose Layer**. These layer use **LeakyReLU as Activation Function** and are followed by a **Batch Normalization Layer**. The network can be viewed as a five block structure with a succession of these layers. In the last block, is used **tanh as Activation Function**, instead of LeakyReLU.

The generator is a fully-convolutional network that inputs a noise vector (latent_dim) to output an image of 3 x 64 x 64. Think of it as a decoder. Feed it a latent vector of

100 dimensions and an upsampled, high-dimensional image of size 3 x 64 x 64 (that represents the generated image) [9].

3.3.2 Discriminator

The Discriminator is always a binary classifier as before, but also here Convolutional Layers are used (in this case, are used classic Convolutional Layers).

The input of the model is an image of dimension 3 x 64 x 64 (3 corresponds to RGB channel and 64 x 64 is pixels size). The Output is always a score between 0 and 1 produced by **Sigmoid** activation function in the output layer.

3.3.3 Training and results

The training on this DCGAN implementation, due to model structure and size (number of parameters to learn) were too expensive in terms of times and computational cost. It took the same time used to train the other model (GAN and cGAN) but in this case the training is done only for 30 epochs instead of 200 and 150 epochs (it tooks 1:30 hour on a *NVIDIA Tesla T4*). But despite the reduced number of training epochs used the results are too good.

In the Table 3.3 is possible to see the generation improvement during the DCGAN's training phase: the images are initially very blurry and some figures does not make sense, instead in the last epoch, they are very sharp and they have the same structure as those of the Anime Face Dataset.



(a) Epoch 1



(b) Epoch 10



(c) Epoch 20



(d) Epoch 30

Table 3.3: Performance improvement during training in terms of generations.

3.4 cDCGAN

Chapter 4

Conclusion

Bibliography

- [1] Afnan Amin Ali. Gan (generative adversarial network). <https://medium.com/swlh/gan-generative-adversarial-network-3706ebfef77e>, 2017. Accessed: 2022-09-22.
- [2] Akshay Cboulwar. The art of convolutional neural network. <https://medium.com/@achoulwar901/the-art-of-convolutional-neural-network-abda56dba55c>, 2019. Accessed: 2022-09-22.
- [3] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks, 2014. URL <https://arxiv.org/abs/1406.2661>.
- [4] Mehdi Mirza and Simon Osindero. Conditional generative adversarial nets. *arXiv preprint arXiv:1411.1784*, 2014.
- [5] Packt. Principles of gans. <https://subscription.packtpub.com/book/big-data-and-business-intelligence/9781788629416/4/ch04lv11sec24/principles-of-gans>, None. Accessed: 2022-09-22.
- [6] Alec Radford, Luke Metz, and Soumith Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. *arXiv preprint arXiv:1511.06434*, 2015.
- [7] Aditya Sharma. Introduction to generative adversarial networks (gans). <https://learnopencv.com/introduction-to-generative-adversarial-networks/>, 2021. Accessed: 2022-09-22.

- [8] Aditya Sharma. Conditional gan (cgan) in pytorch and tensorflow. https://learnopencv.com/conditional-gan-cgan-in-pytorch-and-tensorflow/#CGAN_TensorFlow, 2021. Accessed: 2022-09-22.
- [9] Aditya Sharma. Deep convolutional gan in pytorch and tensorflow. <https://learnopencv.com/deep-convolutional-gan-in-pytorch-and-tensorflow/>, 2021. Accessed: 2022-09-22.